

BRENDAN C. SMITH

Lead Data Scientist

📞 816-522-2998 @ hire@brendansmith.ai 🔗 linkedin.com/in/b-c-s

EXPERIENCE

Lead Data Scientist

09/2024 – 10/2025

Best Egg

Remote

- Created a next-gen customer expansion XGBoost model for the Flexible Rent Platform including alternative data, increasing Gini coefficient by 23%.
- Designed and implemented a series of Metaflow flows (DAGs) to automatically train challenger models for our primary credit risk / underwriting XGBoost model.
- Trained a team of 3 other Data Scientists to contribute flows to the project.
- Worked on a RAG-based chat system to generate answers based on our internal documentation. Also integrated with Snowflake to run SQL queries based on questions.

Senior Data Engineer – Contract

04/2024 – 07/2024

Burns & McDonnell

Kansas City, MO

- Built Airflow DAGs to ingest full-load and incremental data from various SQL databases.
- Transformed records between medallion stages using dbt within the Databricks Medallion Architecture.
- Designed a logging strategy and integrated it with Azure Monitor / Log Analytics.

Graduate Learning Facilitator – Machine Learning

01/2024 – 04/2024

University of Texas at Austin

Remote

- Served as a teaching assistant for a graduate-level Machine Learning course, supporting instruction, grading, and administration for a cohort of approximately 465 students.

Data Scientist

05/2023 – 08/2023

Propense.ai

Remote

- Augmented enterprise knowledge graphs with internal data to cold-start a B2B recommendation system, solving the new-platform data sparsity problem for market insights.
- Identified patterns in sales gaps by analyzing client and sales history data. Presented actionable insights to clients, securing 5 initial contracts for product launch.

Quantitative Investments & Data Science Intern

02/2023 – 05/2023

Nexus Equities

Remote

- Developed a PyTorch computer vision model to estimate the useable land area on outdoor storage facilities using satellite imagery, accelerating approximately 50 investment decisions.

Software Development Engineer – EC2 Core Platform

10/2018 – 01/2021

Amazon Web Services (AWS)

Seattle, WA

- Owned and operated two critical services: one to drain customer instances from unhealthy EC2 hosts, and another to proactively recycle older hosts for re-provisioning. Identified and resolved deadlock conditions, resulting in a \$300k/month reduction in 'unsellable' rate.
- Built a scalable and secure data lake to centralize previously siloed EC2 internal data.
- Used PySpark to ingest real-time data from various AWS data stores such as AWS DynamoDB, RDS (MySQL), S3, and Amazon Athena.
- Deployed ETL pipelines using AWS Glue Jobs, Crawlers, and VPC Elastic Interfaces to extract cross-regional data from 300+ internal AWS production service accounts.

EDUCATION

M.S. in Data Science

University of Texas at Austin

01/2022 – 12/2024

- Created DNNs using PyTorch for the vision system of a racing simulator. Implemented networks for object detection, keypoint estimation, semantic segmentation, multi-action networks, and reinforcement learning for autonomous driving.
- Created neural networks for Natural Language Processing (NLP) applications such as Semantic Parsing & Labelling, Sentiment Analysis, and Language Generation.

GPA | 3.89 / 4.00

B.S. in Computer Science

University of Nebraska – Lincoln

08/2012 – 05/2016

- Student in the honors program, Jeffrey S. Raikes School of Computer Science and Business Management.
- Minor in Business Management and Mathematics.
- Capstone project involved working with local businesses to implement software solutions.

GPA | 3.90 / 4.00

SKILLS

Languages & Frameworks

Python · SQL · R · PySpark · Java · JavaScript
/ TypeScript · C# · PyTorch · NumPy · Pandas
· Polars · Scikit-learn · Matplotlib · Plotly · Seaborn

Machine Learning & MLOps

Deep Learning · Computer Vision · Forecasting · Recommendation Systems · Hyperparameter Tuning · ML Pipelines · Model Monitoring · Metaflow

Generative AI

LLMs · Agentic AI · Tool Use & Function Calling · Fine-tuning · RAG · Prompt Engineering · Embeddings · LangChain

Data Engineering

ETL Pipelines · Data Warehousing · Spark · Kafka · Hive · Hadoop Stack · Databricks

Cloud & Infrastructure

AWS · Azure · Databricks · Docker · REST APIs · FastAPI · SQL & NoSQL Databases ·

- Created stored procedures and QuickSight analytics dashboards, significantly increasing the velocity of insights and business decisions for the TPM team and stakeholders.
- Implemented a capacity forecasting model and integrated it into the proactive re-provisioning workflow, increasing host turnover rate by up to 18% per region.

Software Development Engineer – HDInsight

08/2016 – 10/2018

Microsoft Azure

Redmond, WA

- Served as cross-org liaison for root cause analysis of regressions in Apache Hadoop ecosystem products (Spark, Kafka, etc.), coordinating fixes with upstream Apache engineers.
- Proposed, designed, and deployed time series anomaly detection ML models on Azure, improving alarm triggers and identifying cluster configurations with high customer impact. Reduced average TTD by ~55% and TTR by ~20%.
- Joined a small v-team which refactored the control plane of HDInsight to support the creation of flexible cluster shapes. Closed the feature gap with our primary competitors while increasing service reliability KPIs and decreasing COGS.

Design Studio Software Developer – Decision Science

09/2015 – 05/2016

Hudl

Lincoln, NE

- Collaborated with a team of data scientists to develop a temporal convolutional neural network which automatically detected basketball highlights from user-uploaded video.
- Conducted continuous A/B testing to collect user metrics (engagement, bounce rate, stickiness, and view count) to iterate on model effectiveness.
- Assessed model launch success through user engagement metrics and presented regular reports to stakeholders. Achieved significant growth in active users (DAU / WAU).

Microservices

Data Analysis

Statistical Modeling · Causal Inference · Time Series Analysis · Data Visualization